

Practical Exercise — Lesson 3 (Week 5)

Single-Prediction Explainability Audit with Attention / LIME / SHAP

Time: 75–100 minutes (solo or pairs)

Tools: Google Colab, Python, datasets, transformers, torch, pandas, numpy, matplotlib, lime, shap (optional: bertviz)

What you will submit: a Colab notebook + a 1-page “Explanation Card” (PDF)

Learning outcomes

By the end, you will be able to:

1. Generate **local explanations** for a single prediction using **attention**, **LIME**, and **SHAP**.
 2. Summarize explanations in **plain language** and **visuals** for a non-technical stakeholder.
 3. Run **faithfulness checks** (mask/delete important tokens) to see if explanations predict model behavior.
 4. Document **limitations** (e.g., attention \neq causation; instability; sampling variance) and define next debugging steps.
-

Setup (5–10 min)

1. Open a fresh Colab (CPU is fine).
 2. Install/import:
 3. `!pip -q install datasets transformers torch lime shap matplotlib --upgrade`
 4. `import os, numpy as np, pandas as pd, torch, matplotlib.pyplot as plt`
 5. `from datasets import load_dataset`
 6. `from transformers import AutoTokenizer, AutoModelForSequenceClassification, TextClassificationPipeline`
 7. `os.environ["TOKENIZERS_PARALLELISM"] = "false"`
 8. `SEED = 42`
 9. `np.random.seed(SEED); torch.manual_seed(SEED)`
 10. **Dataset (pick one):**
 - o Quick binary sentiment: glue/sst2 (validation split).
 - o (Optional) Toxicity: civil_comments (small sampled subset).
 11. **Model:** small BERT-like checkpoint for speed, e.g., distilbert-base-uncased-finetuned-sst-2-english.
-

Part A — Pick one instance and lock the baseline (10–15 min)

1. **Load data & model**
 2. `ds = load_dataset("glue", "sst2")["validation"]`
 3. `ckpt = "distilbert-base-uncased-finetuned-sst-2-english"`
 4. `tok = AutoTokenizer.from_pretrained(ckpt, use_fast=True)`
 5. `mdl = AutoModelForSequenceClassification.from_pretrained(ckpt, output_attentions=True)`
 6. `clf = TextClassificationPipeline(model=mdl, tokenizer=tok, return_all_scores=True, truncation=True)`
 7. **Select a single example** whose prediction is confident (score ≥ 0.8) and **save** the raw text and model score.
 8. Record the **baseline prediction** (label + probability). You will use this as the reference for all explanations.
-

Part B — Attention heat-map (last layer, averaged over heads) (15–20 min)

Purpose: Show which tokens receive the most attention from the [CLS] token in the final layer. This is a diagnostic view—not proof of causality.

1. **Run the model with attentions** and decode tokens:
 2. `enc = tok(example_text, return_tensors="pt", truncation=True)`
 3. `with torch.no_grad():`
 4. `out = mdl(**enc)`
 5. `attn = out.attentions[-1].squeeze(0).mean(0) # [heads, seq, seq] -> avg over heads => [seq, seq]`
 6. `cls_to_tok = attn[0].cpu().numpy() # attention from [CLS] (position 0) to all tokens`
 7. `toks = tok.convert_ids_to_tokens(enc["input_ids"][0])`
 8. **Normalize and plot** a bar or heat-map over tokens (skip special tokens). Mark top-k tokens.
 9. **Write a 2–3 sentence narrative:** “The model places high final-layer attention on X, Y, Z, which plausibly aligns/misaligns with the predicted label because ...”
 10. **Limitation note (1 line):** “Attention weights are not causal attributions; they show where the model *looks*, not necessarily which features *drove* the score.”
-

Part C — LIME on the same instance (15–20 min)

Purpose: Obtain perturbation-based word importances with a **local linear surrogate**.

1. **Create a prob-function** that returns `predict_proba` for texts:
 2. `def predict_proba(texts):`
 3. `outs = clf(texts)`
 4. `# outs: list of [{'label':'NEGATIVE', 'score':...},`
`{'label':'POSITIVE','score':...}]`
 5. `# reorder to [neg, pos] and return 2-D array`
 6. `arr = []`
 7. `for o in outs:`
 8. `scores = {d["label"].lower(): d["score"] for d in o}`
 9. `arr.append([scores.get("negative",0.0),`
`scores.get("positive",0.0)])`
 10. `return np.array(arr)`
 11. **Run LIME:**
 12. `from lime.lime_text import LimeTextExplainer`
 13. `class_names = ["negative","positive"]`
 14. `explainer = LimeTextExplainer(class_names=class_names,`
`random_state=SEED)`
 15. `exp = explainer.explain_instance(example_text, predict_proba,`
`num_features=10, labels=[1]) # class 1 = positive`
 16. `exp.show_in_notebook(text=example_text) # or exp.as_list(label=1)`
 17. `lime_weights = exp.as_list(label=1) # list of (token, weight)`
 18. **Interpret (2–3 sentences):** Name the top positive and negative contributors and whether they make sense.
 19. **Limitation note (1 line):** “LIME can be unstable across seeds/perturbations and may split words; treat weights as directional hints.”
-

Part D — SHAP on the same instance (15–20 min)

Purpose: Estimate Shapley values for tokens. This is slower than LIME—keep background small.

1. **Background texts** (10–20 short sentences):
2. `background = [" ".join(ds[i]["sentence"].split()[:10]) for i in`
`np.random.choice(len(ds), 20, replace=False)]`
3. **Define model wrapper** (maps `list[str] → prob` for positive):
4. `def f(texts):`
5. `return predict_proba(texts)[0,1] # positive class probability`
6. **Explainer & values** (kernel approximator for text):
7. `import shap`
8. `masker = shap.maskers.Text(tokenizer=tok)`
9. `explainer = shap.Explainer(f, masker)`
10. `shap_values = explainer([example_text], max_evals=1500, batch_size=20)`
`# keep it light`
11. `shap.plots.text(shap_values[0])`
12. **Interpret (2–3 sentences):** Compare top tokens with LIME. Are they consistent?
13. **Limitation note (1 line):** “Kernel SHAP for text relies on sampling; results vary with background and budget.”

If SHAP is slow or errors: reduce `max_evals` to 500; if still slow, capture `shap_values.data` and compute top-k by absolute value without plotting, or document as a limitation.

Part E — Faithfulness checks (10–15 min)

Goal: Test whether “important” tokens actually **change the prediction** when perturbed.

1. **Deletion test (AOPC-style).**
 - Rank tokens by importance for each method (attention, LIME, SHAP).
 - Iteratively **mask** the top-k tokens (replace with `[MASK]` or a neutral token) and record the positive class probability after each deletion.
 - Plot **probability vs. #tokens removed** curves (three lines: one per method). A steeper drop = more faithful explanation.
 2.

```
def mask_tokens(text, tokens_to_mask):  
3.     words = text.split()  
4.     for i in tokens_to_mask: words[i] = "[MASK]"  
5.     return " ".join(words)  
6. # Build ranked indices per method, then loop k = 1..K
```
 7. **Counterfactual tiny edit.**
 - Change one top positive/negative token (e.g., “excellent”→“average”), re-predict, and record Δ probability.
 - Write a **one-liner**: “Replacing *X* with *Y* reduced the positive probability from 0.91→0.64 (-0.27), consistent with the explanations.”
-

Part F — Synthesis & risks (5–10 min)

1. **Convergence table (top-5 tokens).** Create a 3-column table (Attention / LIME / SHAP) with token and signed weight/indicator. Note overlaps.
 2. **Short narrative (≤ 120 words):**
 - What features are the model using?
 - Do methods agree on the **reason** for the prediction?
 - Any **spurious cues** (punctuation, usernames, identity terms)?
 - One **next step** (data fix, prompt/feature tweak, threshold).
 3. **Caveats (≤ 2 bullets):** One limitation per method you used.
-

What you submit

1. **Notebook (.ipynb)** containing:
 - Baseline prediction; **attention heat-map** or bar chart; **LIME** and **SHAP** visuals/lists.

- **Deletion curve** plot comparing methods; **counterfactual edit** result.
 - Convergence table; short narrative; caveats.
 - 2. **Explanation Card (≤ 1 page, PDF)**
 - **Use-case & instance** (1 sentence, anonymized if needed).
 - **Prediction** (label + prob).
 - **Top tokens** (agreeing across methods, if any).
 - **Faithfulness evidence** (deletion drop, counterfactual Δ prob).
 - **Risk note** (e.g., identity term drives decision).
 - **Action** (one concrete next step) + **how you'll re-measure**.
-

Troubleshooting & tips

- **Tokenization drift:** Use the *same* tokenizer for attention and masking. Skip special tokens ([CLS], [SEP]).
 - **Long texts:** Truncate to ≤ 128 tokens for speed; document the truncation.
 - **LIME stability:** Set `random_state`; if unstable, run twice and note variance.
 - **SHAP speed:** Lower `max_evals` or background size; if still slow, include text output (top tokens) instead of the plot and explain the constraint.
 - **Attention \neq explanation:** Treat it as a **lens**, not ground truth; corroborate with faithfulness checks.
-

Close

You now have a minimal, repeatable **explainability audit** for a single prediction: three views (attention/LIME/SHAP), one **faithfulness test**, and a **plain-English** narrative that drives a concrete engineering decision.